

File Descriptors

Jose L. Muñoz, Oscar Esparza, Juanjo Alins, Jorge Mata
Telematics Engineering
Universitat Politècnica de Catalunya (UPC)

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- 1 File Descrip.
 - Introduction
 - Redirecting Output
 - Redirecting Input
 - Pipes
 - Files in Bash



What is a File Descriptor?

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- An fd is an integer that is used as index of a kernel-resident data structure containing the details of all **open files**.
- Each process has its own “file descriptor table”.
- The same file can have different file descriptors:
 - We can open a file for reading and get one fd.
 - We can open the same file for writing and get another fd.
 - One of the main elements contained in the file descriptor table is the file pointer, which indicates the current position (for r/w) on a file.

Standard fds I

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- There are 3 standard file descriptors which presumably every process should have opened.
- When a process is created using a shell (like Bash), it inherits the 3 standard file descriptors from this shell.
- The `lsOf` command shows us the “list of open files”:

```
$ ps
  PID TTY          TIME CMD
 7988 pts/3        00:00:00 bash
 8839 pts/3        00:00:00 ps
```


```
$ lsOf -a -p 7988 -d0-10
COMMAND  PID    USER      FD  TYPE  DEVICE  NAME
bash     7988   telematics 0r   CHR   136,3   /dev/pts/3
bash     7988   telematics 1w   CHR   136,3   /dev/pts/3
bash     7988   telematics 2w   CHR   136,3   /dev/pts/3
```

- The previous command shows the file descriptors active (up to fd=10) for the process with PID 7988.

Standard fds II

- fd=0 is opened for reading (0r) and fd=1,2 are opened for writing (1w and 2w).

fd (integer value)	Name
0	Standard Input (STDIN)
1	Standard Output (STDOUT)
2	Standard Error (STDERR)



File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- 1 File Descrip.
 - Introduction
 - Redirecting Output**
 - Redirecting Input
 - Pipes
 - Files in Bash



Redirecting Output I

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- The “output redirection” allows to send STDOUT/STDERR to files different from default ones.
- Examples:

```
$ echo Hello, how are you?  
Hello, how are you?
```

```
$ echo Hello, how are you? > file.txt  
$ cat file.txt  
Hello, how are you?
```

```
$ echo Are you ok? >> file.txt  
$ cat file.txt  
Hello, how are you?  
Are you ok?
```

- You can also redirect the standard error (*fd=2*).

Redirecting Output II

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- Example:

```
$ ls -qw  
$ ls -qw 2> error.txt
```

- In general:

- **N>file**. It is used to redirect fd=N to a file. If the file exists, is deleted and overwritten. In case file does not exist, it is created.
- **N>>file**. Similar to the first case but opens file in mode append.
- **& >file**. Redirects STDOUT and STDERR to file.

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- 1 File Descrip.
 - Introduction
 - Redirecting Output
 - Redirecting Input**
 - Pipes
 - Files in Bash



Redirecting Input I

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- Input redirection allows you to specify a file for reading standard input ($fd=0$).
- The format for input redirection is $< file$.
- Example:

```
1 #!/bin/bash
2 # Our third script, using read for fun
3 echo Please, type a sentence and hit ENTER
4 read TEXT
5 echo You typed: $TEXT
```

- To redirect the input:

```
$ echo hello world > file.txt
$ ./thirdscript.sh < file.txt
Please, type a sentence and hit ENTER
You typed: hello world
$
```

Redirecting Input II

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- Let's observe redirections with `lsOf`:

```
1 #!/bin/bash
2 lsOf -a -p $$ -d0-10
3 echo Now, please type a sentence and hit ENTER
4 read TEXT
5 echo You typed: $TEXT
```

```
$ ./thirdscript.sh < file
COMMAND PID USER  FD  TYPE  DEVICE  NAME
script  7728 user   0r   REG   8,6    /home/user/file
script  7728 user   1u  CHR 136,3  /dev/pts/3
script  7728 user   2u  CHR 136,3  /dev/pts/3
Now, please type a sentence and hit ENTER
You typed: hola
```

- The variable `$` contains the PID of the script in execution.

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

Outline

1 File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash



Unnamed pipes I

- Unix based operating systems like Linux offer a unique approach to join two commands on the terminal.
- You can take the output of the first command and use it as input of the second command, this is the concept of pipe or “|”.

```
$ ls | grep x
```

- This type of pipe is called “unnamed pipe” because the pipe exists only inside the kernel and cannot be accessed by processes that created it, in this case, the bash shell.
- All Unix-like systems include a variety of commands to manipulate text outputs.
- We have already seen some of these commands: head, tail, grep and cut.

Unnamed pipes II

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- We also have other commands for text pipelines like:
 - `uniq` which displays or removes repeating lines.
 - `sort` which lists the contents of the file ordered alphabetically or numerically.
 - `wc` which counts lines, words and characters.
 - `find` which searches for files.
- The following example shows a compound command with several pipes:

```
$ cat *.txt | sort | uniq > result.txt
```

- `tee` reads STDIN and writes to a file and also to STDOUT.

```
$ ls | tee output.txt | sort -r
```

Process Substitution I

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- Commands enclosed in parenthesis are run in a “subshell” (cloned shell).
- An application is to redirect the output of several commands:

```
$ (ps ; ls) >commands.out
```

- Process substitution occurs when you put a “<” or “>” in front of the left parenthesis:

```
$ cat <(ls -l)
```

- In the previous example `cat` has a valid file name to read from.
- Similarly, giving “>(commands)” results in bash naming a temporary pipe, which the commands inside the parenthesis read for input:

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

Process Substitution II

```
ls|tee >(grep foo|wc>foo.txt)|grep baz|wc>baz.txt
```


File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- The dash “-” in some commands is useful to indicate that we are going to use stdin or stdout instead of a regular file.
- Example:

```
$ grep linux doc1.txt | cat doc2.txt -
```

- Dash depends on the context (has other uses).
- To give an example, the command `cd -` means go to the previous directory visited (nothing to do with redirection).

Named Pipes

- Another sort of pipe is the “named pipe”, also called FIFO (First In First Out).
- First bytes entered are the first ones to be removed.
- You can create it:

```
$ mkfifo fifo1
```

- To use it:

```
t1$ echo hello how are you? > pipe1
```

```
t2$ cat < pipe
```

- As you will observe, the output of the command run on the first pseudo-terminal shows up on the second pseudo-terminal.

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

1 File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash



Files in Bash I

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- In bash, we can manage a total of 9 file descriptors.
- 0 is the standard input, 1 is the standard output, 2 is the standard error and there are six additional fds that take values from 3 to 9.
- Useful for this example:

```
#!/bin/bash
LOGFILE=/var/log/script.log
comand1 >$LOGFILE
comand2 >$LOGFILE
...
```

Files in Bash II

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

Syntax

```
exec fd> file
exec fd>>file
exec fd< file
exec fd<> file

exec fd1>&fd2

exec fd1<&fd2

exec fd>&-
command >&fd
command 2>&fd
command <&fd
```

Meaning

open file for writing and assign `fd`.
open file for appending and assign `fd`.
open file for reading and assign `fd`.
open file for reading/writing and assign `fd`.

open `fd1` for writing. From this moment `fd1` and `fd2` are the same.

open `fd1` for reading. From this moment `fd1` and `fd2` are the same.

close `fd`.
write stdout to `fd`.
write stderr to `fd`.
read stdin from `fd`.

Files in Bash III

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- Let us illustrate the use of `exec` for reading files by an example:

```
#!/bin/bash
LOGFILE=/var/log/script.log
exec 1>$LOGFILE
cmd1
cmd2
...
```

```
$ exec 3>&1          # create a new fd that points to
                   # the same place as stdout (fd=1)
$ exec 1>script.log  # fd=1 is now redirected to the log file
$ cmd1              # the output of cmd1 will go into the log file.
$ cmd2              # the output of cmd2 will go into the log file.
$ exec 1>&3          # now fd=1 points the same file as fd=3
$ cat script.log    # Now we see the output in our terminal
$ lsof -a -p $$ -d0-3 # we have four fd open for this bash
```

Files in Bash IV

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- In another example let us assign the file file.log to fd=3 and use this file descriptor number for writing, finally we close the fd.

```
$ exec 3>file.log
$ lsof -a -p $$ -d0,1,2,3
COMMAND  PID  USER  FD   TYPE DEVICE SIZE  NODE NAME
bash     3443 student 0u    CHR 136,35        37 /dev/pts/35
bash     3443 student 1u    CHR 136,35        37 /dev/pts/35
bash     3443 student 2u    CHR 136,35        37 /dev/pts/35
bash     3443 student 3u    REG   3,1      0 86956 /home/student/file.log
$ echo hello >&3
$ cat file.log
hello
$ exec 3>&-
```

Files in Bash V

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- Open a file in read mode works similarly:

```
#!/bin/bash
exec 4<&0
exec <restaurant.txt
while read score type phone
do
    echo $score,$type,$phone
done
exec 0<&4
exec 4>&-
```


File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- Opening in read/write mode

```
$ echo 1234567890 > numbers.txt
$ exec 3<> numbers.txt
$ read -n 4 <&3           # read 4 characters (moves the pointer)
$ echo -n . >&3           # write a dot (without LF)
$ exec 3>&-               # close fd=3
$ cat numbers.txt
1234.67890
```

Files in Bash VII

File Descrip.

Introduction

Redirecting Output

Redirecting Input

Pipes

Files in Bash

- You can also do redirections using different *fd* but only for a single command line (not for all commands executed with the bash).

```
#!/bin/bash
exec 3>student.log          # Open fd=3 (for bash)
echo "This goes to student.log" 1>&3  # redirects stdout to student.log
                                   # only for this command line

echo "This goes to stdout"
exec 1>&3                    # redirects stdout to student.log
                                   # for the rest of the commands

echo "This also goes to student.log"
echo "and this sentence, too"
```

- Children processes inherit the opened *fd* of their parent process.
- The children can close an *fd* if it is not going to be used.